

MICROSERVICES LAB

Reference link for spring-boot application

<https://blog.scottlogic.com/2019/10/31/building-microservices-with-spring-boot.html>

EXERCISES ON KUBERNETIES with solution

Create a simple deployment of the given app with name of your choice and 3 replicas of pods. Check the status of pod by sending request. App should be accessed from outside the cluster. Use Port Forwarding to Access Applications in a Cluster.

dep.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: usn-nginx-deployment
  labels:
    app: usn-nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: usn-nginx
  template:
    metadata:
      labels:
        app: usn-nginx
    spec:
      containers:
        - name: nginx
          image: 172.1.14.168:5001/nginx
          ports:
            - containerPort: 80
```

Command to deploy:

```
kubectl apply -f dep.yaml
```

Command to check pods

```
kubectl get pods | grep 'usn' \ \ type your usn
```

Command to expose

```
kubectl expose deployment usn-nginx-deployment --type=NodePort --name=usn-nginx-service
```

To get exposed port

```
kubectl get svc | grep 'usn' \ \ type your usn
http://172.1.14.168:<nodeport>
```

```
kubectl port-forward deployment/usn-nginx-deployment newport:<nodeport>
```

Demonstrate the updation of image in live container in a pod using command line as well as by updating yaml files

dep.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: usn-nginx-deployment
  labels:
    app: usn-nginx
spec:
  replicas: 3
  selector:
    matchLabels:
      app: usn-nginx
  template:
    metadata:
      labels:
        app: usn-nginx
    spec:
      containers:
        - name: nginx
          image: 172.1.14.168:5001/nginx
          imagePullPolicy: "Always"
          ports:
            - containerPort: 80
```

Command to deploy:

```
kubectl apply -f dep.yaml
```

Command to expose

```
kubectl expose deployment usn-nginx-deployment --type=NodePort --name=usn-nginx-service
```

Command to update image:

```
kubectl set image deployment/usn-nginx-deployment nginx=newImage
```

To check the updated name:

```
kubectl describe deploy usn-nginx-deployment | grep Image:
```

Perform the following.

- a. Create 3 pods with names nginx1, nginx2,nginx3. All of them should have the label app=v1 Show all labels of the pods.
- b. Get only the 'app=v2' pods.
- c. Remove the 'app' label from the pods we created before

```
kubectl run usn-nginx1 --image=nginx --restart=Never --labels=app=usn-v1
```

```
kubectl run usn-nginx2 --image=nginx --restart=Never --labels=app=usn-v1
```

```
kubectl run usn-nginx3 --image=nginx --restart=Never --labels=app=usn-v1
```

```
kubectl get po --show-labels
```

```
kubectl get po -l app=usn-v2
```

```
kubectl label po nginx1 nginx2 nginx3 app-
```

Create a Pod with ubuntu image and a command to echo "YOUR_NAME" which overrides the default CMD/ENTRYPOINT of the image.

dep_ubuntu_pod1.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu
  labels:
    app: ubuntu
spec:
  containers:
  - name: ubuntu
    image: 172.1.14.168:5001/ubuntu
    command: ["/bin/bash"]
    args: ["-c", "echo MSRIT"]
```

```
kubectl apply -f dep_ubuntu_pod1.yaml
```

```
kubectl logs ubuntu
```

```
kubectl exec --stdin --tty ubuntu -- /bin/bash
```

```
kubectl delete pod ubuntu
```

Create a Pod that runs one container. The configuration file for the Pod defines a command and arguments by using environment variables:

dep_ubuntu_pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu
  labels:
    app: ubuntu
spec:
  containers:
    - name: ubuntu
      image: 172.1.14.168:5001/ubuntu
      env:
        - name: MESSAGE
          value: "MSRIT"
      command: ["/bin/echo"]
      args: ["$(MESSAGE)"]
```

```
kubectl apply -f dep_ubuntu_pod.yaml
kubectl logs ubuntu
kubectl delete pod ubuntu
```